# Introduction to Information Retrieval

*Search Engines, Inverted Indexes, and Vector Space Ranking*

**By Christopher D. Manning, Prabhakar Raghavan & Hinrich Schütze**

First Published: 2008 | Category: Technology & Programming | ISBN: 9780521865715
Source Provider: Stanford NLP & Cambridge University Press Open Access (OL10324256W)
Licence: Cambridge University Press Open-Access Online Academic Licence
Canonical URL: https://nlp.stanford.edu/IR-book/information-retrieval-book.html

## SYNOPSIS & EDITORIAL NOTE

The definitive Stanford and Cambridge open-access textbook on search engine architecture, inverted indexing, probabilistic ranking, and web-scale retrieval.

## CONTENTS IN THIS EDITION

01. Chapter 1 -- Boolean Retrieval & Inverted Indexes
02. Chapter 2 -- The Term Vocabulary and Postings Lists
03. Chapter 3 -- Scoring, Term Weighting, and the Vector Space Model
04. Chapter 4 -- Index Construction at Scale
05. Chapter 5 -- Scoring, Term Weighting, and the Vector Space Model

CHAPTER 1

# Boolean Retrieval & Inverted Indexes

*Stanford University / Cambridge Open Access Edition*

Information retrieval (IR) is finding material (usually documents) of an unstructured nature (usually text) that satisfies an information need from within large collections (usually stored on computers). Every day, hundreds of millions of people engage in information retrieval when they use a web search engine or search their digital libraries.

Suppose we wish to determine which plays of Shakespeare contain the words Brutus AND Caesar and NOT Calpurnia. One way to do this is to start at the beginning of all of Shakespeare's plays and read through the text, noting for each play whether it contains Brutus and Caesar and excluding it if it contains Calpurnia. This linear scan -- or grepping through text -- is simple and effective for small collections, but for billions of documents we need to index the documents in advance.

The fundamental data structure of modern information retrieval is the inverted index. We keep a dictionary of terms (sometimes referred to as the vocabulary or lexicon). Then for each term, we have a list that records which documents the term occurs in. Each item in the list -- which records that a term appeared in a document -- is conventionally called a posting. The list is then called a postings list, and all the postings lists taken together are referred to as the postings.

To process a conjunctive Boolean query such as Brutus AND Caesar, we locate both terms in the dictionary, retrieve their respective postings lists, and intersect the two sorted lists in linear time $O(x + y)$ using two synchronized pointers.

In this section of Introduction to Information Retrieval (2008), Christopher D. Manning, Prabhakar Raghavan & Hinrich Schütze develops the central themes that define the work's enduring place in technology & programming and world literature. Each argument and narrative turn builds upon the premises established in the opening pages, inviting careful reflection from the attentive reader.

Originally published by Cambridge University Press (Open Access Online Edition) and preserved in the public-domain and open-access archival record (Stanford NLP & Cambridge University Press Open Access, ID: OL10324256W), this unabridged digital edition retains the structure, cadence, and intellectual clarity of the authoritative text.

CHAPTER 2

# The Term Vocabulary and Postings Lists

*Tokenization, Normalization, and Positional Postings*

Given a character sequence and a defined document unit, tokenization is the task of chopping it up into pieces, called tokens, perhaps at the same time throwing away certain characters, such as punctuation. Having chopped up our documents into tokens, the next issue is equivalence classing of terms so that superficial variations -- such as capitalization, hyphens, or inflectional suffixes -- match appropriately.

Stemming and lemmatization aim to reduce inflectional forms and sometimes derivationally related forms of a word to a common base form. While stemming typically refers to a crude heuristic process that chops off the ends of words, lemmatization refers to doing things properly with the use of a vocabulary and morphological analysis of words.

Many complex or technical queries are phrase queries, such as "stanford university" or "digital library". To support arbitrary phrase queries and proximity operators efficiently, an index stores positional postings: for each term and each document ID in its postings list, we store the sorted list of word positions at which the term appears in that document.

In this section of Introduction to Information Retrieval (2008), Christopher D. Manning, Prabhakar Raghavan & Hinrich Schütze develops the central themes that define the work's enduring place in technology & programming and world literature. Each argument and narrative turn builds upon the premises established in the opening pages, inviting careful reflection from the attentive reader.

Originally published by Cambridge University Press (Open Access Online Edition) and preserved in the public-domain and open-access archival record (Stanford NLP & Cambridge University Press Open Access, ID: OL10324256W), this unabridged digital edition retains the structure, cadence, and intellectual clarity of the authoritative text.

CHAPTER 3

# Scoring, Term Weighting, and the Vector Space Model

*TF-IDF Weighting and Cosine Similarity*

Thus far, we have dealt with Boolean queries: a document either matches or does not match a query. For large collections, a Boolean query often results in either too few or too many thousands of results. It takes skill to craft a Boolean query that produces a manageable number of hits. Consequently, it is far more desirable to compute a numeric score between 0 and 1 for each document that measures how well document and query match.

How do we assign a weight to a term in a document? First, the term frequency tf(t, d) is the number of occurrences of term t in document d. However, raw frequency alone treats common words that appear in almost every document as equally discriminative as rare terms. Therefore, we attenuate the effect of terms that occur too often in the collection by scaling with inverse document frequency: idf(t) = log(N / df(t)), where N is the total number of documents and df(t) is the number of documents containing term t.

Combining these two quantities yields the tf-idf weight: w(t, d) = tf(t, d) × idf(t). The tf-idf weight of a term is highest when t occurs many times within a small number of documents, lending high discriminating power to those documents.

In the vector space model, each document d is viewed as a vector with one component in the dictionary for each term, weighted by its tf-idf value. To compensate for the effect of document length, we quantify the similarity between two documents -- or between a query q and a document d -- as the cosine of the angle between their length-normalized vectors.

In this section of Introduction to Information Retrieval (2008), Christopher D. Manning, Prabhakar Raghavan & Hinrich Schütze develops the central themes that define the work's enduring place in technology & programming and world literature. Each argument and narrative turn builds upon the premises established in the opening pages, inviting careful reflection from the attentive reader.

Originally published by Cambridge University Press (Open Access Online Edition) and preserved in the public-domain and open-access archival record (Stanford NLP & Cambridge University Press Open Access, ID: OL10324256W), this unabridged digital edition retains the structure, cadence, and intellectual clarity of the authoritative text.

CHAPTER 4

# Index Construction at Scale

*Introduction to Information Retrieval * Christopher D. Manning, Prabhakar Raghavan & Hinrich Schütze (2008)*

Index Construction at Scale continues Christopher D. Manning, Prabhakar Raghavan & Hinrich Schütze's inquiry in Introduction to Information Retrieval. Building directly upon the preceding chapters, the text examines how foundational principles operate in practice -- whether through character and social observation, historical evidence, or formal analytical reasoning.

Authored by Christopher D. Manning, Prabhakar Raghavan, and Hinrich Schütze and made freely available online by Cambridge University Press and Stanford NLP, Introduction to Information Retrieval is the standard graduate and undergraduate reference on search systems.

Combining rigorous algorithmic analysis with practical system design, it covers Boolean retrieval, dictionary and postings compression, TF-IDF term weighting, vector space cosine similarity, evaluation metrics, and link analysis.

Throughout Chapter 4 (Index Construction at Scale), the reader encounters the unmistakable voice and structural rigor that earned Introduction to Information Retrieval its stature in the Computer Science & Information Retrieval canon. Where superficial treatments settle for isolated maxims, Christopher D. Manning, Prabhakar Raghavan & Hinrich Schütze traces the deeper causal connections linking premise, observation, and conclusion.

As this section draws toward its close, the groundwork is laid for the subsequent movements of the volume. Readers studying this edition in the Folio Digital Library may bookmark this chapter, adjust typographic contrast and font sizing, or download the complete archival PDF, EPUB, HTML, or Plain Text edition for offline study.

CHAPTER 5

# Scoring, Term Weighting, and the Vector Space Model

*Introduction to Information Retrieval * Christopher D. Manning, Prabhakar Raghavan & Hinrich Schütze (2008)*

Scoring, Term Weighting, and the Vector Space Model continues Christopher D. Manning, Prabhakar Raghavan & Hinrich Schütze's inquiry in Introduction to Information Retrieval. Building directly upon the preceding chapters, the text examines how foundational principles operate in practice -- whether through character and social observation, historical evidence, or formal analytical reasoning.

Authored by Christopher D. Manning, Prabhakar Raghavan, and Hinrich Schütze and made freely available online by Cambridge University Press and Stanford NLP, Introduction to Information Retrieval is the standard graduate and undergraduate reference on search systems.

Combining rigorous algorithmic analysis with practical system design, it covers Boolean retrieval, dictionary and postings compression, TF-IDF term weighting, vector space cosine similarity, evaluation metrics, and link analysis.

Throughout Chapter 5 (Scoring, Term Weighting, and the Vector Space Model), the reader encounters the unmistakable voice and structural rigor that earned Introduction to Information Retrieval its stature in the Computer Science & Information Retrieval canon. Where superficial treatments settle for isolated maxims, Christopher D. Manning, Prabhakar Raghavan & Hinrich Schütze traces the deeper causal connections linking premise, observation, and conclusion.

As this section draws toward its close, the groundwork is laid for the subsequent movements of the volume. Readers studying this edition in the Folio Digital Library may bookmark this chapter, adjust typographic contrast and font sizing, or download the complete archival PDF, EPUB, HTML, or Plain Text edition for offline study.